

Organizations That Learn with AI

Hillary Fan

GitHub reading selection G1 · 20 September 2026

A guide to this selection

People solve difficult problems at work, but the next team often has to solve them again. AI can make an individual task faster without helping an organization remember what it learned. These were among the questions Hillary Fan pursued in the essays that led toward EIP.

Five short selections follow that development: writing to understand a problem, finding connections, building organizational capability, examining supervision, and noticing the work people quietly do to keep AI useful. Author excerpts retain the supplied wording; editorial notes connect them to the later framework.

This is a historical reading companion, not the current EIP specification or a report of validated outcomes. Current definitions belong to the separate EIP public repository. Selection and source details appear at the end.

1. Writing as a way to investigate

Source: content unit 01, “Why I started writing,” P26.

Author excerpt — three selected paragraphs; surrounding and intervening material omitted:

I was trying to answer one question that kept following me throughout my career: How can organizations collaborate with increasingly capable AI without sacrificing trust, accountability, professional identity, or data sovereignty?

Writing became my way of thinking. Every article started as a question—not a conclusion.

But every answer led to another question. Who owns the knowledge generated during AI collaboration?

Editorial note: The starting point is a working problem: making useful AI participation compatible with responsibilities that organizations already carry. This explains why the later framework concerns decision rights and reusable knowledge, rather than model access alone.

2. Seeing connections across separate questions

Source: content unit 02, “The architecture taking shape,” P27.

Author excerpt — one opening paragraph and one continuation paragraph; intervening material omitted:

After writing 23 articles, I decided to step back and read them again. What surprised me most was not the individual ideas. It was the pattern. Looking back today, I no longer see 23 separate articles. I see one evolving narrative about how organizations might adapt to increasingly capable AI.

One final reflection. I didn't write these 23 articles because I believed I had already found the answer. I wrote them because writing became the way I tested my own thinking. Every conversation challenged an assumption. Every disagreement forced me to clarify another idea. Looking back, perhaps these articles aren't a finished framework. They are simply an open record of how my thinking evolved. If that journey resonates with someone facing similar organizational questions, then sharing it was worthwhile.

Editorial note: “23 articles” belongs to this earlier retrospective, not the size of the later collected edition. The full source also explores credentialing, transformation economics, and speculative agency. Those explorations are not all current EIP commitments. The current public framework separates its core operating model from optional research directions and does not establish portable credentials or self-funding results.

3. From local productivity to organizational capability

Source: content unit 38, “From AI Adoption to Enterprise Capability,” P40.

Author excerpt — two consecutive paragraphs:

Over the past few years, I kept noticing the same pattern. Organizations were becoming more productive. Yet they didn't seem to become significantly more capable. Frontline experts solved complex problems every day. AI helped people work faster. New ideas emerged constantly. Yet somehow, valuable operational experience rarely became organizational assets.

Instead, knowledge remained within individuals, teams repeatedly solved similar problems, and organizations continued to depend on a handful of experienced experts. That observation eventually led me to a different question:

Editorial note: These are the author's observations and problem framing, not a measured comparison of enterprise performance. The question connects to EIP's present distinction between temporary work, required records, and approved reusable capability. Retaining

everything would not answer it: a saved conversation is not automatically reliable institutional knowledge.

The practical question is whether useful experience can be selected, checked, scoped, and reused without removing accountable ownership. EIP proposes an operating model for testing that possibility.

4. Asking what human supervision contributes

Source: content unit 44, “What Human-in-the-Loop Is Not Answering,” P46.

Author excerpt — three consecutive paragraphs:

Much of the current discussion on Human-in-the-Loop focuses on how to reduce review burden while maintaining safety, quality, and control.

That is important — but I do not think it is the whole question.

A deeper unresolved possibility is this:

Author excerpt — the immediately following question:

What if human supervision is not only an operating requirement, but also a formation stage?

Editorial note: The source asks whether intensive early cognitive support could contribute to capability that later survives a reduction in that support. It does not demonstrate a declining supervision curve. The present v2.1 distinction is important: evidence about reduced cognitive support does not remove an authority checkpoint or grant greater permission to act.

This question connects to Black Sheep's investigation of transferable control structures. The connection is conceptual. Neither a reported mature trajectory nor an essay establishes cross-user generalization or a reliable mechanism.

5. Making the hidden work visible

Source: content unit 51, “6-5 | Human-Absorbed Effort: Who Completes the AI?”, P53.

Author excerpt — two consecutive opening paragraphs:

AI output may appear complete because people repeatedly repair prompts, verify facts, correct formats, resolve agent conflicts, restore missing context, take over failed steps, manage affected stakeholders, and translate results into explanations an organization can accept.

This labor is dispersed across roles, messages, and tickets. It rarely appears in model cost or beside the completion rate. I call it human-absorbed effort: human work that

disappears from the measurement of AI performance precisely because people successfully absorb it.

Editorial note: This is the author's proposed framing of a measurement problem. This selection omits the original essay's survey figures, external evidence discussion, and rhetorical closing; it makes no fresh quantitative claim. The practical question remains whether reported AI performance includes the human repair and coordination needed to obtain it.

That concern appears in several current projects: EIP asks about cognitive sustainability; Health Companion proposes reducing information reconstruction; Black Sheep investigates context repair and handoff. These are related questions, not mutually validating results.

Where the work stands

The collection makes the author's development path more visible. It does not independently settle authorship or originality, and it is not evidence that the proposed systems have been implemented. The supplied edition credits the author's concepts and direction while disclosing AI assistance with translation, source work, editing, and production. This GitHub adaptation was also AI-assisted and owner-authorized.

Earlier claims remain historical even when readable or persuasive. The current public EIP baseline governs present framework descriptions. Its complete volumes, controlled methods, and private development records are not published in this selection.

Further reading and sources

The excerpts reproduce selected paragraphs from the supplied revised Word edition, not complete essays. Surrounding and, where indicated, intervening paragraphs are omitted. Editorial notes are new. Source links identify posts listed in the supplied collection; their live contents and publication dates were not independently verified for this edition. The current EIP v2.1 public baseline is a reviewed integration draft awaiting bounded organizational testing.

[EIP current public framework](#)

[EIP v2.1 public baseline](#)

[My Champion Labs project guide](#)

[Source links, original-byte hashes, and editorial record](#)

© 2026 Hillary Fan. The repository NOTICE applies. No additional open-source or implementation license is granted.